


Pearls Karachi AQI Predictor

Presenter: Muhammad Hamza Zeeshan
Data Science, 10Pearls Shine Internship Cohort 7



2. Project Overview

This project focuses on building a **self-sustaining machine learning ecosystem** that predicts air quality for Karachi utilizing a **100% serverless stack**.

Objective

The primary goal is to architect an **end-to-end MLOps pipeline** that:

- **Automates Data Ingestion:** Fetches live pollutant and weather data every hour via API.
- **Ensures Model Freshness:** Implements daily retraining to adapt to new atmospheric patterns.
- **Provides Reliable Forecasts:** Delivers a rolling **72-hour AQI prediction** with high accuracy

3. Data Acquisition

For this project, I integrated the **Open-Weather API** as the primary data provider. It was selected because:

- Up to 40 years back (hourly/daily) Historical Data
- High-resolution **Karachi-specific coordinates**
- Ability to provide atmospheric weather and air quality variables in a single, unified response.

Historical Data Backfilling

To quickly build a training set, I created a **backfilling script** (`backfill_data.py`) to programmatically request **historical data** for Karachi. This yielded **~4,500 records (~7 months of hourly data)** before live data collection began.

4. Exploratory Data Analysis (EDA)


I processed ~4,500 records to create a high-quality baseline dataset for Karachi's air quality forecasting model. This involved transforming raw API data, identifying patterns, and ensuring data integrity.

- **Profiling & Statistical Summary:** Analysis confirmed that AQI levels frequently fluctuate between "Moderate" and "Poor," allowing me to define normal ranges for pollutants like PM2.5 and CO before training.
- **Cleaning & Imputation:** I removed duplicates and handled gaps using imputation techniques.
- **Correlation & Key Findings:** The data revealed a strong direct link between PM2.5 concentrations and the final AQI score.


A significant discovery was the **diurnal cycle**, which showed pollution peaks during Karachi's morning and evening rush hours. These insights confirmed that time-of-day is a mandatory feature for the model to achieve high-accuracy 72-hour predictions.

5. Data Visualization

Time-series line charts showed pollutant levels, especially PM2.5 and CO, spiking bimodally around 9:00 AM and 5:00 PM, confirming the heavy impact of rush-hour vehicular emissions.



Scatter plots and heatmaps revealed that atmospheric conditions influence pollution. A key finding was the inverse correlation between wind speed and AQI: higher wind speeds disperse pollutants, while stagnant, humid air leads to "Hazardous" AQI levels.



6. Feature Engineering

I engineered a specialized set of **17 features** to capture the complexity of Karachi's atmosphere:

- **Temporal Features:** Extracted **hour, day of the week, and month** to model cyclical traffic and seasonal patterns.
- **Lagged Features:** Created **1-hour lags** for AQI and key pollutants (e.g., aqi_lag_1h), which provide the model with essential "memory" of recent conditions.
- **Dynamic Features:** Calculated the **AQI change rate** to identify how quickly pollution levels are rising or falling.

Feature Store Integration (Hopsworks)

I integrated the **Hopsworks Feature Store** to manage features at scale. It acts as a centralized "data warehouse," with the **Feature Pipeline** automatically inserting hourly updates, guaranteeing both training and inference pipelines consistently access the same, up-to-date data.

karachi_aqi_fg version 1

	Created at	Last updated at	Mode	Online							
	2026-01-05 16:37:47	2026-02-16 16:00:50	Stream	false							
features	17	commits	100	feature views	1	time travel	hudi	data validation	none	monitoring config	0

Hourly AQI data for Karachi with time-based features and 1-hour lags

7. Model Development

Multiple architectures were benchmarked to find the most effective model:

- **Ridge Regression:** Baseline for linear relationships.
- **Random Forest Regressor:** Selected for superior handling of non-linear patterns.
- **Neural Networks:** DL model to capture long-term atmospheric dependencies.

Training & Hyperparameter Tuning

The models were trained on a historical dataset of **~4,500 records**. To prevent overfitting and ensure the model generalizes well to new data, I implemented:

- **Regularization:** Used in neural architectures to maintain lean, effective weight distributions.
- **Early Stopping:** Automated the training cutoff to stop once the validation error plateaued.
- **Grid Search:** Fine-tuned hyperparameters to minimize the **Mean Absolute Error (MAE)**.

Model Registry (Version Control)

To manage the evolution of the AI, I utilized the **Hopsworks Model Registry** allowing me to:

- **Track Experiments:** Logged over **25 model versions**, comparing their accuracy scores side-by-side.
- **Champion/Challenger Strategy:** Only promoted the best-performing model.
- **Metadata Storage:** Saved training logs and feature importance maps alongside the model file for full transparency.

Version Metrics & Deployments					
Version displayed					
	version	mae	r2	Deployments	
		min 0.0087	max 0.1963	min 0.8296	max 0.9959
	29	0.1251		0.9168	not deployed
	28	0.1462		0.9226	not deployed
	27	0.1537		0.9106	not deployed
	26	0.1526		0.9135	not deployed
	25	0.1566		0.9074	not deployed

8. Model Evaluation

8.1 Performance Metrics

I evaluated the model using industry-standard regression metrics to quantify its predictive power:

- **R² Score:** This high "Coefficient of Determination" indicates that the model explains over 83% of the variance in Karachi's air quality data.
- **Mean Absolute Error (MAE):** On average, the model's predictions are within a very small margin of the actual recorded AQI values.



8.2 Prediction Results

The AI accurately predicts bimodal rush-hour peaks across a 72-hour horizon. Validated on unseen data with high precision, the system is production-ready for delivering reliable early warnings of hazardous air quality.

8.3 Feature Importance Analysis

SHAP analysis validated that the model identifies the correct drivers for Karachi's air quality:

- **Primary Drivers:** `aqi_lag_1h` and **PM2.5** concentrations are the most influential factors.
- **Secondary Signals:** **Temporal features (hour)** confirm the impact of daily traffic cycles on pollution.



9. Automation & Pipelines

All workflows
Showing runs from all workflows

460 workflow runs

- Hourly AQI Data Fetch (Scheduled, main)

Data Fetch

The hourly **Feature Pipeline** pulls raw API data and processes 17 engineered features directly into the Hopsworks Feature Store.

- **Feature Ingestion:** Updates the model's "memory" with the latest air quality trends.
- **Automation:** Triggered every 60 minutes to ensure data freshness for Karachi.

Model Training

The **Training Pipeline** automates model selection and registration.

- **Chronological Validation:** Utilizes a professional 80/20 time-series split.
- **Benchmarking:** Trains three distinct architectures (**Ridge**, **Random Forest**, and **Neural Network**) with aggressive regularization (high alpha, shallow depth, 40% dropout) to avoid over-fitting on noise.
- **Model Registry:** Automatically selects the "Winner" based on the lowest MAE, exports performance metrics to `model_info.json`, and registers the model to the Hopsworks Registry.

Inference



The **Inference Pipeline** converts stored features into actionable 72-hour forecasts. It retrieves the best model from the registry to compute future AQI levels.

- **Batch Forecasting:** Generates a rolling 3-day outlook in a single automated run.
- **Data Delivery:** Pushes results as a CSV back to the repository to act as a live data feed for the Streamlit dashboard.

10. Deployment: Interactive Dashboard

I developed a high-performance Streamlit application to bridge live environmental data with actionable predictive insights for Karachi's citizens.

- **Sober UI Design:** Data readability with a 1–5 color-coded AQI status indicator
- **Performance Optimization:** Implemented `st.cache_data` to ensure near-instant load
- **Alerts:** Translates raw AQI scores into red-zone warnings and specific advice.



11. Technology Stack

- **Core Logic:** Python 3.11, Pandas, and NumPy for all pipeline data processing.
- **Machine Learning:** Scikit-Learn (Random Forest, Ridge) for modeling and SHAP for interpretability.
- **MLOps Infrastructure:** Hopsworks as the central Feature Store and Model Registry hub.
- **Orchestration:** GitHub Actions for automated hourly/daily YAML-based workflows.
- **Frontend:** Streamlit for the UI and Open-Weather API for real-time meteorological streams.

12. Challenges & Solutions

I resolved several technical bottlenecks to ensure the stability and accuracy of the end-to-end system:

- **Missing Libraries:** GitHub Runners failed because they couldn't find Scikit-Learn.
Solution: Updated the .yml configuration to ensure all necessary tools are installed automatically before the script runs.
- **Data Type Conflicts:** Hopsworks rejected data because Python's numbers were too large (64-bit vs 32-bit).
Solution: Added a step to the pipeline to force all data into the correct format before sending it to the feature store.
- **Access Denied:** The dashboard could not load the forecast file. **Solution:** Adjusted repository settings to ensure the frontend had the correct permissions to fetch the results.
- **Schedule Delays:** Automated tasks didn't always start on time at first.
Solution: Monitored the system behavior for several days to ensure the triggers became stable and consistent.
- **Static Predictions:** The model was "memorizing" data instead of predicting, causing it to output the same numbers.
Solution: Fixed the overfitting by adding constraints to the model's complexity, forcing it to learn actual patterns.

13. Conclusion: Achievements & Key Learnings

My internship at **10Pearls** was a transformative experience, applying MLOps to a production-grade project within a professional, distributed engineering workflow.

- **Professional Mentorship:** Received expert guidance from architecture design to final deployment.
- **Technical Networking:** Collaborated with a cohort of like-minded computer science students.
- **Serverless Mastery:** Built a scalable system using **Serverless Architecture** without managing physical infrastructure.
- **MLOps Proficiency:** Implemented a **Central Feature Store** and **Model Registry** for professional data/model versioning.
- **Skill Integration:** Synthesized **Python, Data Science, and DevOps** skills into a cohesive solution.